

---

# Optimization Algorithms, teaching notes

IllusionCraft

*Prepared By: Abhijit Singh Jowhari*

## 1. Gradient Descent

### Introduction

Gradient Descent is one of the most fundamental optimization algorithms used in machine learning. It iteratively adjusts parameters to minimize a given loss function.

### Mathematical Formulation

Given a loss function  $L(\theta)$ , where  $\theta$  represents the parameters of the model, the goal is to find the parameters  $\theta$  that minimize  $L$ .

1. **Loss Function:**  $L(\theta)$  2. **Gradient:** The gradient of the loss function  $\nabla_{\theta}L(\theta)$  is a vector of partial derivatives of  $L$  with respect to  $\theta$ . 3. **Update Rule:**

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta}L(\theta_t)$$

where  $\eta$  is the learning rate.

### Detailed Steps

1. **Initialization:** Initialize the parameters  $\theta$  randomly. 2. **Compute Gradient:** Calculate the gradient of the loss function at the current parameters. 3. **Update Parameters:** Adjust the parameters in the direction of the negative gradient. 4. **Repeat:** Continue steps 2 and 3 until convergence (when the change in the loss function is below a threshold or after a fixed number of iterations).

### Convergence and Learning Rate

- **Learning Rate ( $\eta$ ):** A small learning rate may lead to slow convergence, while a large learning rate can cause divergence. Often, the learning rate is chosen through experimentation or using techniques like learning rate schedules. - **Convergence:** The algorithm converges when  $\|\nabla_{\theta}L(\theta)\|$  is close to zero.

### Advantages and Disadvantages

- **Advantages:** - Simplicity and ease of implementation. - Suitable for convex optimization problems. - **Disadvantages:** - Can be slow for large datasets. - Sensitive to the choice of learning rate. - Can get stuck in local minima for non-convex problems.

---

## 2. Stochastic Gradient Descent (SGD)

### Introduction

Stochastic Gradient Descent is a variant of Gradient Descent where the gradient is estimated using a single or a small subset of data points, rather than the entire dataset.

### Mathematical Formulation

#### 1. Loss Function:

$$L(\theta) = \frac{1}{N} \sum_{i=1}^N L_i(\theta)$$

where  $N$  is the number of training samples and  $L_i$  is the loss for the  $i$ -th sample.

**Stochastic Gradient:** The gradient is estimated using a single sample or a mini-batch:

$$\nabla_{\theta} L(\theta) \approx \nabla_{\theta} L_i(\theta)$$

#### 3. Update Rule:

$$\theta_{t+1} = \theta_t - \eta \nabla_{\theta} L_i(\theta_t)$$

### Detailed Steps

1. **Initialization:** Initialize the parameters  $\theta$  randomly. 2. **Shuffling:** Randomly shuffle the training data. 3. **Iterate Over Data:** For each sample  $i$  in the training data: - Compute the gradient of the loss function for the sample. - Update the parameters using the computed gradient. 4. **Repeat:** Continue the process for a fixed number of epochs or until convergence.

### Advantages and Disadvantages

- **Advantages:** - Faster iterations since gradients are computed on fewer data points. - Can escape local minima due to its stochastic nature. - Often reaches a good enough solution faster than batch gradient descent. - **Disadvantages:** - Higher variance in the parameter updates can lead to a less stable convergence path. - Requires more iterations to converge compared to batch gradient descent. - The noisy updates may cause the algorithm to never converge to the exact minimum.

## 3. Adam (Adaptive Moment Estimation)

### Introduction

Adam is an advanced optimization algorithm that combines the benefits of two other extensions of stochastic gradient descent: Adaptive Gradient Algorithm (AdaGrad) and Root Mean Square Propagation (RMSProp).

---

## Mathematical Formulation

Adam maintains two moving averages of the gradient: the first moment (mean) and the second moment (uncentered variance).

1. **Initialize Parameters:** -  $m_0 = 0$  (first moment vector) -  $v_0 = 0$  (second moment vector) -  $t = 0$  (time step) 2. **Hyperparameters:** -  $\alpha$ : Learning rate -  $\beta_1$ : Decay rate for the first moment (typically 0.9) -  $\beta_2$ : Decay rate for the second moment (typically 0.999) -  $\epsilon$ : A small constant to prevent division by zero (typically  $10^{-8}$ )

3. **Update Rule:** - Increment time step:  $t = t + 1$  - Compute gradient:  $g_t = \nabla_{\theta} L(\theta_t)$  - Update biased first moment estimate:

$$m_t = \beta_1 m_{t-1} + (1 - \beta_1) g_t$$

- Update biased second moment estimate:

$$v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$$

- Correct bias in first moment:

$$\hat{m}_t = \frac{m_t}{1 - \beta_1^t}$$

- Correct bias in second moment:

$$\hat{v}_t = \frac{v_t}{1 - \beta_2^t}$$

- Update parameters:

$$\theta_{t+1} = \theta_t - \alpha \frac{\hat{m}_t}{\sqrt{\hat{v}_t} + \epsilon}$$

## Detailed Steps

1. **Initialization:** Initialize the parameters  $\theta$ ,  $m_0$ ,  $v_0$ , and set  $t = 0$ . 2. **Iterate:** For each step: - Increment  $t$ . - Compute the gradient  $g_t$ . - Update the biased first and second moment estimates. - Compute bias-corrected estimates. - Update the parameters. 3. **Repeat:** Continue until convergence or for a predetermined number of iterations.

## Advantages and Disadvantages

- **Advantages:** - **Adaptive Learning Rate:** Adjusts the learning rate for each parameter individually. - **Bias Correction:** Provides bias-corrected first and second moment estimates, leading to more accurate updates. - **Efficiency:** Combines the advantages of AdaGrad (good for sparse gradients) and RMSProp (good for non-stationary objectives). - **Robustness:** Well-suited for large-scale problems and high-dimensional parameter spaces. - **Disadvantages:** - **Complexity:** More complex to implement and understand compared to simpler algorithms like SGD. - **Resource Intensive:** Requires more memory to store moment estimates.

## Practical Considerations

### Choosing Hyperparameters

- **Gradient Descent:** Learning rate ( $\eta$ ). - **SGD:** Learning rate ( $\eta$ ), batch size. - **Adam:** Learning rate ( $\alpha$ ),  $\beta_1$ ,  $\beta_2$ ,  $\epsilon$ .

---

Hyperparameters significantly influence the performance of optimization algorithms. Techniques such as grid search, random search, and Bayesian optimization can help in tuning these parameters.

## Implementation in Libraries

Modern machine learning libraries like TensorFlow, PyTorch, and Keras provide built-in implementations of these optimization algorithms. For example:

- In **TensorFlow**:

```
optimizer = tf.optimizers.Adam(learning_rate=0.001)
```

- In **PyTorch**:

```
optimizer = torch.optim.Adam(model.parameters(), lr=0.001)
```

- In **Keras**:

```
optimizer = keras.optimizers.Adam(learning_rate=0.001)
```

## Real-life Examples

1. **Gradient Descent in Linear Regression:** - Used to find the best-fit line by minimizing the mean squared error between the predicted and actual values.
2. **SGD in Image Classification:** - Commonly used in training deep neural networks for tasks like image recognition, where the dataset is large and using the entire dataset for each update is computationally expensive.
3. **Adam in Natural Language Processing (NLP):** - Frequently used in training transformers and other deep learning models for tasks such as language translation, where adaptive learning rates and robust convergence properties are crucial.

## Monitoring and Debugging

- **Track Loss:** Monitor the loss function value over iterations to ensure the algorithm is converging. - **Learning Rate Adjustments:** If the loss is not decreasing, try adjusting the learning rate. - **Visualization:** Plot the loss function and parameter values to understand the optimization path and detect issues like vanishing gradients or exploding gradients.

## Practical Tips

- **Initialization:** Proper initialization of parameters can prevent issues like slow convergence or getting stuck in local minima. Techniques such as Xavier or He initialization are commonly used. - **Batch Size:** In SGD, choosing an appropriate batch size is critical. A small batch size introduces noise but can help escape local minima, while a large batch

---

size provides more stable updates. - **Learning Rate Schedules:** Dynamic learning rate adjustments (e.g., reducing the learning rate when the improvement in loss is minimal) can improve convergence.

**End of Document**